

GEAR: A General Experience-Driven Approach for Cost-Efficient Software Engineering Agents

Anonymous ACL submission

Abstract

Software engineering (SE) agents powered by large language models are increasingly adopted in practice, yet they often incur substantial monetary cost. We introduce GEAR, a general experience-driven approach that reduces the cost of SE agents while preserving task performance. GEAR extracts structured experience from prior issue-resolution executions and leverages it to guide patch generation and selection, enabling early termination of unproductive iterations. We evaluate GEAR on the SWE-bench Verified benchmark across three representative agents. GEAR consistently reduces total cost by 19%–55% (32% on average), with negligible loss in resolution rate (at most 0.2%). These efficiency gains are achieved, on average, by identifying early-termination opportunities for 11% of issues and reducing API calls, input tokens, and output tokens by 21%, 30%, and 25%, respectively. We release the code, prompts, and data at <https://github.com/EffiSEAgent/GEAR>.

1 Introduction

Large language models (LLMs) are reshaping modern software development, with software engineering (SE) agents emerging as one of their most impactful applications (Liu et al., 2025). Given an issue describing a bug or a feature request along with the corresponding code repository, an SE agent aims to navigate the codebase, localize relevant code, and generate a patch (i.e., modify the code) to resolve the issue, with correctness validated by the associated tests (Xia et al., 2025).

To achieve this, a variety of SE agents have been developed. For example, Agentless (Xia et al., 2025) follows a fixed, expert-designed workflow without autonomous planning. In contrast, Mini-SWE-Agent (Yang et al., 2024) and Trae Agent (Gao et al., 2025) are autonomous, capable of interacting with the environment, using tools, and planning multi-step actions. Mini-SWE-Agent

focuses on patch generation, whereas Trae Agent performs both patch generation and selection.

Despite the remarkable capabilities of existing SE agents, cost efficiency remains a key concern for practical adoption. In a recent Stack Overflow survey (StackOverflow, 2025), 53% of software engineers reported that the cost of using agents is a barrier for them. This monetary cost is also associated with environmental impact due to energy consumption (Ren et al., 2024) and other resource usage (Jegham et al., 2025). Reducing cost inefficiency is therefore important for both economic and environmental reasons.

This cost inefficiency arises largely because SE agents typically resolve issues through multi-round iterations (Gao and Peng, 2026), such as repeated tool invocations, reflecting the inherent complexity of software development. The cost grows super-linearly with each iteration: in every round, the entire conversation history, including prompts, tool calls, and outputs, is fed back as context, causing computational costs to escalate rapidly.

To address this challenge, we propose GEAR, a general experience-driven approach for cost reduction of SE agents while preserving task performance. The key insight behind GEAR is that, much like an experienced software engineer, an agent equipped with relevant experience does not need to explore many iterations to resolve an issue.

GEAR realizes this through two complementary mechanisms. First, it captures structured experience from historical issue-resolution activities, encoding information about the issue, the agent’s execution trajectory, and the outcomes of prior attempts. Second, when tackling a new issue, GEAR retrieves relevant experience to guide the agent’s behavior, enabling early termination of redundant iterations during both patch generation and selection. By leveraging past experience, GEAR reduces unnecessary computational cost while maintaining or even improving the likelihood of success-

084
085
086
087
088
089
090
091
092
093
094
095
096
097
098
099
100
101
fully resolving the issue. Importantly, GEAR is a highly general optimization approach that can be integrated seamlessly into diverse agents without requiring fundamental redesigns.

We evaluate GEAR on the widely adopted SWE-bench Verified benchmark (OpenAI, 2025) using three representative SE agents, namely Agentless (Xia et al., 2025), Mini-SWE-Agent (Yang et al., 2024), and Trae Agent (Gao et al., 2025), across different LLM backends. The results show that GEAR consistently reduces total cost by 19.3%–55.1% (31.8% on average), with negligible loss in resolution rate (at most 0.2%). These efficiency gains are achieved by identifying early-termination opportunities for 8.6%–14.0% of issues (11.3% on average), as well as by reducing API calls, input tokens, and output tokens by 20.8%, 29.9%, and 25.1% on average, respectively.

2 Related Work

103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
SE Agents. SWE-bench (Jimenez et al., 2024a) and its human-validated variant SWE-bench Verified (OpenAI, 2025) have become standard benchmarks for evaluating automated resolution of real-world SE issues. Motivated by these benchmarks, recent work has focused on agent-based approaches that equip LLMs with tool use and iterative execution to autonomously explore codebases and generate patches (Liu et al., 2025). Representative examples include Mini-SWE-Agent (Yang et al., 2024), which relies on shell-based interaction for iterative code navigation and editing, and Trae Agent (Gao et al., 2025), which combines parallel patch generation with a selector agent that leverages static analysis and execution traces. While these agents achieve strong performance, previous work (Fan et al., 2025) has shown that such performance is accompanied by substantial token consumption.

121
122
123
124
125
126
127
128
129
130
131
132
133
Agent Efficiency Optimization. Previous work on improving agent efficiency can be grouped into three categories. First, prompt-level approaches reduce input overhead through compression (Jiang et al., 2023) or Retrieval-Augmented Generation (RAG) (Shinn et al., 2023). Second, agent re-architectures lower cost via strategies such as plan caching (Zhang et al., 2025), codified prompting (Yang et al., 2025), or multi-agent delegation (Gandhi et al., 2024). Third, reasoning constraints control runtime resource use, e.g., per-turn token budgeting (Han et al., 2025), static safeguards (Jimenez et al., 2024a), or dynamic turn-

control (Gao and Peng, 2026). While effective at reducing cost, these approaches typically degrade task performance, highlighting the challenge of achieving both efficiency and performance.

Agent Memory. Recent agents incorporate memory mechanisms to improve task performance, such as shared message pools in MetaGPT (Hong et al., 2024) and hierarchical memory designs in Generative Agents (Park et al., 2023) and MemoryBank (Zhong et al., 2024). More broadly, RAG (Lewis et al., 2020) extends agent context via external memory. The experience used in GEAR can be viewed as a form of agent memory. However, previous work primarily uses memory to improve task performance and often incurs additional computational overhead due to frequent retrieval and context expansion. For example, MemoryBank (Zhong et al., 2024) requires frequent memory retrieval operations, leading to increased cost. In contrast, GEAR leverages structured experience specifically to reduce agent cost.

3 Methodology

155
156
157
158
159
160
161
Figure 1 presents an overview of GEAR, a general experience-driven approach for improving the cost efficiency of SE agents. It collects experience from past agent executions of issue resolution and leverages this experience to guide patch generation and selection in a cost-efficient manner.

162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
GEAR consists of three components: experience generation, experience-driven patch generation, and experience-driven patch selection. Experience generation records agent execution trajectories, i.e., sequences of tool calls and their outcomes, together with resolution results from historical issues, and summarizes them into an experience base. For a new issue, the SE agent retrieves relevant experience objects during patch generation to guide tool usage and enable early stopping of unproductive generation. Once a patch and its execution trajectory are passed to the patch selection stage, GEAR evaluates them against retrieved experience to both guide efficient selection and determine whether further patch generation or iteration is necessary.

3.1 Experience Generation

178
179
180
181
182
The experience generation component aims to distill an agent’s historical issue resolution experience into an experience base that can effectively support the resolution of new issues. During issue

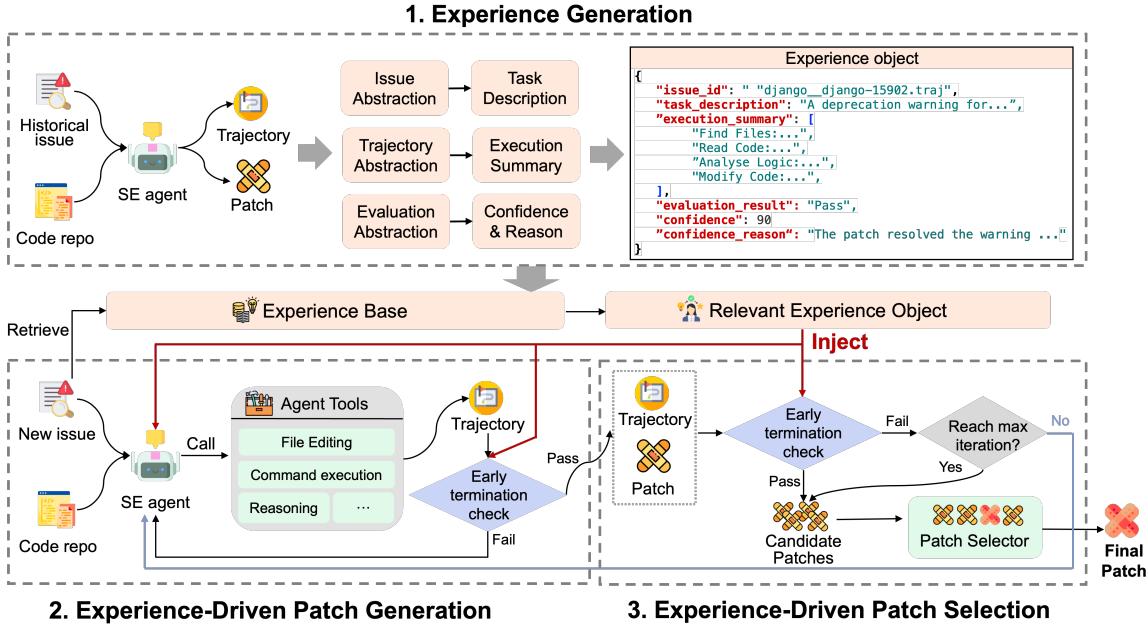


Figure 1: Overview of GEAR.

resolution, agents produce complex execution trajectories consisting of tool calls and their execution outcomes, which serve as the primary source of experience. However, storing raw trajectories introduces substantial noise and token overhead, while overly aggressive simplification may discard actionable signals (e.g. location and reasoning). We therefore design a balanced experience representation that preserves decision-relevant information while remaining compact and efficient.

Experience Representation. Figure 1 illustrates an example schema of an experience object (a complete example is provided in Appendix A). Each experience object captures an agent’s experience from resolving a single issue and is tagged with the corresponding issue ID. The remaining fields provide a high-level abstraction of the issue resolution task, the resolution process, and the resolution outcome, enabling the experience to generalize across related issues. This abstraction is performed using the agent’s backend LLM. In the following, we describe each field in detail.

Issue abstraction: Real-world issue descriptions are often lengthy and detailed (Jimenez et al., 2024b), containing background information, discussion, and auxiliary context that is not directly relevant to resolution. For efficiency, we store a concise abstraction of each issue in the task_description field, capturing its core intent and technical requirements. This abstraction serves as a stable semantic anchor for matching and re-

trieving relevant experience for future issues.

Trajectory abstraction: Rather than recording every tool invocation and execution detail, we compress each execution trajectory into a sequence of key resolution steps, storing in the execution_summary field. The abstraction process filters out non-essential intermediate actions and retains only critical steps that influence the final outcome, significantly reducing token usage compared to storing full trajectories.

Evaluation abstraction: Issue resolution typically includes a patch evaluation phase, which validates whether a generated patch resolves the issue by executing tests. In GEAR, we retain only high-quality experience, namely resolution trajectories that lead to a successfully validated patch. From an efficiency perspective, unsuccessful attempts primarily reflect exploratory trial-and-error and provide limited reusable value, while successful resolutions capture actionable signals that can directly inform early stopping and decision-making. Accordingly, the evaluation_result field is set to pass for all stored experiences.

To further distill high-level and reusable knowledge from these resolution experiences, we introduce two additional fields: confidence, which provides a quantitative estimate of the reliability of the resolution process, and confidence_reason, which offers a structured retrospective assessment of the generated patch along predefined criteria, including completeness, quality, relevance, and

245 supporting evidence. Together, these fields form
246 lightweight quantitative and qualitative abstrac-
247 tions of experience. By compressing raw execution
248 traces into structured, high-value knowledge, they
249 can serve as an efficiency driver that enables agents
250 to bypass redundant trial-and-error and terminate
251 unproductive iterations earlier.

252 **Experience Retrieval.** Experiences collected from
253 historical issue resolutions are stored in an experi-
254 ence base. To enable their effective reuse in future
255 issue resolution, the experience base supports expe-
256 rience retrieval. Given a new issue, we retrieve rel-
257 evant experience objects by applying a commonly
258 used text similarity approach (TF-IDF (Salton and
259 Buckley, 1988)) to match the current issue descrip-
260 tion against the task_descriptions stored in the
261 experience base.

262 The retrieval strategy is designed to balance both
263 task performance and efficiency. To preserve task
264 performance, we mitigate context pollution, where
265 irrelevant experiences waste tokens and distract
266 the agent, by enforcing a strict similarity threshold
267 τ_{sim} . To ensure efficiency, we retrieve only the top-
268 1 experience whose similarity to the current issue
269 exceeds this threshold. This ensures that the agent
270 is provided with a highly relevant, compact, and
271 high-quality reference with minimal overhead.

272 3.2 Experience-Driven Patch Generation

273 During the patch generation phase, SE agents iter-
274 atively produce candidate patches for a given issue
275 by interacting with the code repository and execu-
276 tion environment. This process typically involves
277 invoking a suite of tools, e.g., file editing tools for
278 modifying source code, command execution tools
279 for running tests and capturing outputs or errors,
280 and reasoning tools for structured analysis and de-
281 cision making. The sequence of tool invocations
282 and their outcomes is recorded as a trajectory.

283 In existing SE agents, patch generation typically
284 terminates only when the backend LLM explic-
285 itly call termination or when a predefined iteration
286 budget is exhausted (Gao and Peng, 2026). Such
287 coarse-grained termination criteria often lead to un-
288 necessary iterations and excessive cost, especially
289 when the agent continues refining patches that are
290 already likely to resolve the issue.

291 GEAR addresses this through experience guid-
292 ance and early termination. For a new issue, GEAR
293 first retrieves the relevant experience object and
294 conditions the agent’s patch generation on both
295 the issue description and the retrieved experience,

296 improving generation effectiveness. More impor-
297 tantly, GEAR introduces an experience-driven early
298 termination mechanism based on intermediate mile-
299 stones during patch generation.

300 Specifically, GEAR treats both code modifica-
301 tion and test execution as milestones. After each
302 milestone, the agent estimates a confidence score
303 indicating whether the current modifications are
304 sufficient to resolve the issue, informed by retrieved
305 experience. If the confidence score exceeds a
306 threshold τ^{gen} , patch generation is terminated early;
307 otherwise, the agent continues iterating. Patch gen-
308 eration is also terminated when the maximum iter-
309 ation budget of the underlying SE agent is reached.

310 This dual-milestone design reflects the practical
311 observation that confidence in a patch may arise
312 either immediately after code modifications, based
313 on structural and semantic alignment with past suc-
314 cessful fixes, or after observing dynamic feedback
315 from test executions. By enabling early termina-
316 tion at both stages, GEAR avoids redundant iterations
317 while preserving patch quality.

3.3 Experience-Driven Patch Selection

318 SE agents commonly rely on repeated patch genera-
319 tion to exploit the stochasticity of LLMs. In prac-
320 tice, an agent generates a predetermined number
321 of candidate patches, bounded by a maximum iter-
322 ation limit, and forwards them to a patch selector.
323 The selector evaluates candidates using reproduc-
324 tion tests, which verify whether the original issue is
325 resolved, and regression tests, which check that pre-
326 viously passing tests remain valid. Given the high
327 execution cost of SE agents, such fixed repetition
328 is often inefficient. Simple issues may be resolved
329 with only a single generation, while difficult or
330 ill-posed issues may exhibit early signs of failure,
331 making further attempts unlikely to succeed.

332 To address this inefficiency, we propose an
333 experience-driven patch selection mechanism that
334 dynamically controls the number of generated
335 patches. After each patch is generated, GEAR
336 evaluates its potential using three inputs: the patch
337 itself, the corresponding execution trajectory, and
338 relevant experience retrieved from past issue reso-
339 lutions. Based on these inputs, the backend LLM
340 is prompted to produce a confidence score indi-
341 cating the likelihood that the patch will lead to a
342 successful resolution. To improve reliability and
343 interpretability, the LLM is also required to provide
344 a brief rationale for its confidence assessment.

345 Early termination decisions are made by com-

paring the confidence score against two thresholds, $\tau_{\text{upper}}^{\text{sel}}$ and $\tau_{\text{lower}}^{\text{sel}}$. If the score exceeds $\tau_{\text{upper}}^{\text{sel}}$, the patch is deemed sufficiently reliable, and further patch generation is terminated to eliminate redundancy. Conversely, if the score falls below $\tau_{\text{lower}}^{\text{sel}}$, the issue is considered unlikely to be resolved by the current agent configuration, and additional resource-intensive iterations are halted to avoid waste. In both cases, the accumulated patches are passed to the patch selector to produce the final output. If the confidence score lies between the two thresholds and the maximum iteration limit has not been reached, the agent proceeds to generate additional patches. Otherwise, once the iteration limit is reached, all accumulated patches are forwarded to the patch selector.

Unlike the early termination check during patch generation, which avoids unnecessary cost within the generation of a single patch, experience-driven patch selection prevents wasteful generation of additional patches. Together, these two complementary early termination checks enable GEAR to substantially improve the cost efficiency of SE agents without compromising task performance.

4 Evaluation Setup

This section describes the experimental setup used to evaluate GEAR.

4.1 Research Questions (RQs)

We aim to evaluate GEAR by answering the following RQs.

RQ1 (Effectiveness of GEAR): What is the impact of GEAR on task performance and cost?

RQ2 (Comparison with baselines): How does GEAR compare with baseline methods?

RQ3 (Ablation study): How do experience injection and early termination individually affect GEAR’s effectiveness?

4.2 Datasets

We evaluate GEAR on SWE-bench Verified (OpenAI, 2025), a human-validated benchmark released by OpenAI for reliably assessing AI models’ ability to resolve SE issues. It has been the most widely adopted dataset for evaluating SE agents (SWE-bench Team, 2025). The dataset consists of 500 real-world tasks collected from Python repositories on GitHub. Each task is derived from a resolved GitHub issue and is accompanied by a set of unit tests used to validate candidate patches.

For each task, agents are provided with the original GitHub issue description as the problem statement and granted access to the corresponding code repository. Based on this information, agents are required to modify the repository files to fix the issue. The unit tests are withheld from the agents and are used solely for evaluation.

To generate experience for GEAR, we additionally use SWE-bench Lite (Jimenez et al., 2024a), a commonly used auxiliary dataset that contains 300 tasks following the same format as SWE-bench Verified. To prevent data leakage, we remove any tasks that overlap with SWE-bench Verified. After deduplication, 207 tasks remain, which are used exclusively for experience generation.

4.3 SE Agents

We evaluate GEAR on three representative, state-of-the-art open-source SE agents: Agentless (Xia et al., 2025), Mini-SWE-Agent (Yang et al., 2024), and Trae Agent (Gao et al., 2025). They cover three different paradigms of SE agent design. Agentless follows a fixed, expert-designed workflow without autonomous planning or complex tool use, while Mini-SWE-Agent and Trae Agent are autonomous agents capable of tool use, environment interaction, and multi-step planning. Mini-SWE-Agent directly generates a single patch, whereas Trae Agent performs both patch generation and selection.

Because Agentless follows a predefined, fixed workflow for patch generation and does not support autonomous tool invocation, GEAR cannot be applied to its patch generation phase and is instead applied only to patch selection. In contrast, Mini-SWE-Agent does not include a patch selection phase; therefore, GEAR is applied solely during patch generation. For Trae Agent, which supports both patch generation and selection, GEAR is applied to both phases.

For each agent, we implement two variants using GPT-5-mini and DeepSeek-V3.2 as the LLM backends. These two models are widely adopted representatives of closed-source and open-source LLMs, respectively. All agents are implemented with their default hyper-parameters.

4.4 Baseline Methods

We compare GEAR with two baseline methods:

- *Turn-control:* Gao and Peng (Gao and Peng, 2026) recently propose a turn-control approach to reduce the cost of SE agents, where a turn denotes a single tool invocation. The agent is

Agent	% Resolved	# API Calls	# Input Tokens	# Output Tokens	Total Cost
GPT-5-mini					
Agentless with GEAR Change	33.2%	4,500	35,919,315	1,861,379	\$13.77
	41.0%	3,310	17,306,929	912,076	\$6.18
	+7.8%	-26.4%	-51.8%	-51.0%	-55.1%
Mini-SWE-Agent with GEAR Change	55.6%	7,250	87,818,230	3,775,009	\$16.07
	56.6%	6,679	75,766,697	3,635,735	\$12.96
	+1.0%	-7.9%	-13.7%	-3.7%	-19.4%
Traq Agent with GEAR Change	66.2%	58,607	1,802,021,581	23,187,391	\$161.79
	66.2%	41,086	1,254,073,114	16,686,359	\$116.18
	0.0%	-29.9%	-30.4%	-28.0%	-28.2%
DeepSeek-V3.2					
Agentless with GEAR Change	42.6%	4,500	36,996,894	1,898,606	\$11.21
	49.8%	3,351	25,183,691	1,234,093	\$7.60
	+7.2%	-25.5%	-31.9%	-35.0%	-32.2%
Mini-SWE-Agent with GEAR Change	66.4%	33,896	614,853,823	6,696,703	\$52.20
	67.0%	31,054	531,220,952	6,402,107	\$42.15
	+0.6%	-8.4%	-13.6%	-4.4%	-19.3%
Traq Agent with GEAR Change	70.2%	151,038	4,979,212,134	39,770,353	\$170.01
	70.0%	111,039	3,102,450,888	28,550,220	\$107.56
	-0.2%	-26.5%	-37.7%	-28.2%	-36.7%
Average Change	+2.7%	-20.8%	-29.9%	-25.1%	-31.8%

Table 1: (RQ1) Comparison of task performance and cost of SE agents with and without GEAR. Changes in % Resolved are reported as absolute differences, while changes in other metrics are computed as relative changes.

constrained by an initial turn limit and granted a one-time extension if the limit is reached before completion. Following the original work, we set the initial limit to the 25th percentile of turn counts observed on SWE-bench Lite, and extend it to the 50th percentile if needed.

- *Naive-GEAR*: RAG has been widely adopted to improve SE agents by retrieving relevant context to augment generation (Tao et al., 2025). To isolate the effect of experience abstraction in GEAR, we implement a variant, termed Naive-GEAR, which retrieves raw execution trajectories and generated patches as context, instead of the structured experience objects used in GEAR. All other components and settings are kept identical to those of GEAR.

4.5 Evaluation Metrics

We evaluate GEAR in terms of both task performance and cost efficiency.

Performance Metric. We use the resolution rate (denoted as *% Resolved*) as the metric, defined as the percentage of tasks for which an agent produces a patch that passes all tests (OpenAI, 2025).

Efficiency Metrics. We comprehensively measure cost efficiency from multiple perspectives:

- *# API Calls*: The total number of LLM API invoca-

cations across the benchmark.

- *# Input/Output Tokens*: The total number of input and output tokens consumed by LLMs across the benchmark.
- *Total Cost*: The total monetary cost incurred across the benchmark, measured in U.S. dollars.

4.6 Implementation Details

To tune the hyper-parameters of GEAR, including τ_{sim} , τ_{gen} , $\tau_{\text{upper}}^{\text{sel}}$, and $\tau_{\text{lower}}^{\text{sel}}$, we randomly select 50 issues from the SWE-bench dataset (Jimenez et al., 2024a) that are not included in SWE-bench Verified or SWE-bench Lite to serve as validation data. Detailed hyper-parameter tuning procedures are provided in Appendix B.

5 Results

This section answers our research questions with evaluation results.

5.1 RQ1: Effectiveness of GEAR

This RQ evaluates the effectiveness of GEAR in terms of its impact on both task performance and cost. Table 1 compares performance and cost with and without GEAR across six agent configurations, instantiated by Agentless, Mini-SWE-Agent, and Traq Agent with two LLM backends (i.e., GPT-5-

Agent	GPT-5-mini	DeepSeek-V3.2
Agentless	14.0%	11.8%
Mini-SWE-Agent	10.6%	8.6%
Trae Agent	12.4%	10.2%
Average		11.3%

Table 2: (RQ1) Early termination rates of different agents with GEAR.

mini and DeepSeek-V3.2). Overall, GEAR substantially reduces cost while preserving and, in some cases, improving task performance.

Impact on cost. GEAR reduces total cost by 19.3%–55.1% across the six configurations, with an average reduction of 31.8%. This cost saving is consistently observed across different agent paradigms and LLM backends. For Mini-SWE-Agent, the cost reduction remains significant but is smaller than that of the other two agents. This is because GEAR affects only the cost of generating a single patch for Mini-SWE-Agent, whereas for the other agents it can reduce cost by avoiding repeated patch generation-selection cycles, resulting in larger overall savings.

Cost reductions are also consistent across multiple dimensions. Specifically, GEAR reduces the number of API calls by 7.9%–29.9%, input tokens by 13.6%–51.8%, and output tokens by 3.7%–51.0%, with average reductions of 20.8%, 29.9%, and 25.1%, respectively.

Examining cost from another perspective, we consider the early termination rate of issues. Table 2 presents the results: with GEAR, the agents achieve early termination for an average of 11.3% of issues, ranging from 8.6% to 14.0% across different agents.

Impact on performance. Across the six agent configurations, GEAR maintains task performance with negligible loss. Only a minor decrease of 0.2% is observed for Trae Agent with DeepSeek-V3.2; however, McNemar’s test (McNemar, 1947) indicates that this difference is not statistically significant ($p = 0.460$). For Trae Agent with GPT-5-mini and Mini-SWE-Agent with both GPT-5-mini and DeepSeek-V3.2, resolution rates remain effectively unchanged after applying GEAR. Notably, for Agentless, GEAR increases the resolution rate by 7.8% and 7.2% with GPT-5-mini and DeepSeek-V3.2, respectively. This improvement is likely because Agentless exhibits relatively lower resolution rates, allowing the experience-driven guidance in

GEAR to have a more substantial impact on task performance.

Ans. to RQ1: GEAR substantially reduces total cost by 19.3%–55.1% (average 31.8%) across the evaluated SE agents, with negligible loss in resolution rate (at most 0.2%). In particular, it reduces API calls by 7.9%–29.9% (average 20.8%), input tokens by 13.6%–51.8% (average 29.9%), and output tokens by 3.7%–51.0% (average 25.1%). It also achieves early termination for 8.6% to 14.0% of issues (average 11.3%).

5.2 RQ2: Comparison with Baselines

This RQ compares GEAR with two representative baseline methods: Turn-control and Naive-GEAR. As described in Section 4.3, Agentless follows a predefined, fixed workflow for patch generation and does not support autonomous tool invocation; therefore, turn-control, which explicitly limits tool invocation turns, cannot be applied to Agentless.

To save space, Table 3 reports the average effects of the three methods on task performance and cost across the evaluated SE agents, while detailed results for each agent are provided in Table 8.

Comparison with Turn-control. While Turn-control achieves a larger reduction in cost, it does so at the expense of substantial task performance loss. On average, Turn-control reduces total cost by 41.1%, compared to 31.8% achieved by GEAR; however, it also decreases resolution rate by 10.7%, whereas GEAR improves resolution rate by 2.7% on average. This cost-performance pattern is consistently observed across all agent configurations evaluated, as shown in Table 8.

These results indicate that GEAR strikes a superior balance between cost and task performance. This advantage stems from GEAR’s ability to dynamically decide whether to terminate iterations early based on the agent’s current execution status and previously collected experience, allowing it to avoid unnecessary iterations without prematurely stopping productive ones.

Comparison with Naive-GEAR. Overall, GEAR outperforms Naive-GEAR in both cost reduction and task performance preservation. While Naive-GEAR also leverages previous execution trajectories to enable early termination and reduce cost, these raw trajectories are larger and less structured than the experience objects used by GEAR, result-

Method	% Resolved	# API Calls	# Input Tokens	# Output Tokens	Total Cost
Turn-control	-10.7%	-33.9%	-45.3%	-31.8%	-41.4%
Naive-GEAR	-3.1%	-16.6%	-22.7%	-20.8%	-26.0%
GEAR	+2.7%	-20.8%	-29.9%	-25.1%	-31.8%

Table 3: (RQ2) Average impact of GEAR and baseline methods on task performance and cost across the evaluated SE agents.

ing in higher overhead. Thus, GEAR achieves greater cost reduction. In addition, the unstructured nature of the raw trajectories provides less effective guidance for patch generation and selection, whereas GEAR’s structured experience objects better support decision-making, leading to improved task performance.

Ans. to RQ2: Overall, GEAR achieves a better balance between cost reduction and task performance: while Turn-control reduces cost more aggressively, it incurs a substantial performance loss (10.7% on average), and GEAR outperforms Naive-GEAR in both cost savings and performance preservation.

5.3 RQ3: Ablation Study

This RQ investigates the individual contributions of experience injection and early termination in GEAR. To this end, we conduct an ablation study with two variants: (1) disabling experience injection while retaining early termination, and (2) disabling early termination in patch generation and selection while retaining experience injection. Due to budget constraints, this experiment is performed only on Trae Agent, which is the most advanced agent among those evaluated. We use GPT-5-mini as the LLM backend because it incurs lower inference cost than DeepSeek-V3.2.

As shown in Table 1, Trae Agent augmented with GEAR achieves the same resolution rate as the original Trae Agent, while reducing total cost, API calls, input tokens, and output tokens by 28.2%, 29.9%, 30.4%, and 28.0%, respectively.

Impact of experience injection. When experience injection is removed, the resolution rate drops from 66.2% to 55.8%, although API calls, input tokens, and output tokens, total cost are reduced by 41.9%, 41.7%, 40.7%, and 58.9%, respectively. This performance degradation is expected. Without experience injection, early termination relies solely on local or heuristic signals and may prematurely terminate promising patch generation or

selection processes. Experience injection provides historical guidance that helps distinguish unproductive iterations from those likely to yield valid patches. Removing this guidance increases the risk of terminating effective search paths, leading to lower task performance, even though cost is further reduced due to the absence of experience-related token overhead.

Impact of early termination. In contrast, when early termination is disabled, the resolution rate increases from 66.2% to 66.6%, while API calls, input tokens, output tokens, and total cost increase by 9.2%, 7.3%, 8.0% and 3.1%, respectively. This behavior is also expected: disabling early termination allows the agent to explore more iterations during patch generation and selection, increasing the likelihood of finding valid patches. However, this comes at the expense of higher cost, as experience-related token overhead is incurred without the compensating benefit of early stopping.

Ans. to RQ3: Removing experience injection leads to a substantial drop in resolution rate, despite further reducing cost. In contrast, disabling early termination slightly improves resolution rate, but leads to increased cost.

6 Conclusion

We present GEAR, a general experience-driven framework that reduces the cost of SE agents while preserving task performance. By leveraging structured experience from prior executions and applying experience-guided early termination during patch generation and selection, GEAR mitigates cost inefficiencies caused by redundant iterations in current SE agents. Our evaluation across multiple agent paradigms and LLM backends demonstrates consistent cost reductions (31.8% on average) with negligible loss on resolution rate (at most 0.2%). These gains stem from reducing unnecessary API calls, limiting token usage, and identifying opportunities for early termination.

645 Limitations

646 While GEAR demonstrates significant cost reductions
647 and maintains task performance, this paper
648 has several limitations.

649 *(1) Evaluation scope.* Consistent with prior SE
650 agent studies (Gao et al., 2025; Gao and Peng,
651 2026), our evaluation is limited to SWE-bench
652 Verified, a human-validated benchmark released by
653 OpenAI for reliably assessing AI models’ ability
654 to resolve SE issues. In future work, we plan to
655 extend the evaluation to additional benchmarks as
656 more reliable datasets become available. If feasi-
657 ble, we also aim to evaluate GEAR in industrial
658 settings.

659 *(2) Dependence on historical data.* As an
660 experience-driven approach, GEAR relies on his-
661 torical data to construct its experience base. For
662 entirely novel issues or domains with sparse his-
663 torical data, i.e., the “cold start” problem, GEAR
664 may fail to retrieve relevant experience. This lim-
665 itation is common to RAG-based approaches and
666 motivates further exploration of methods that gen-
667 eralize better in low-data scenarios.

668 References

669 Zhiyu Fan, Kirill Vasilevski, Dayi Lin, Boyuan Chen,
670 Yihao Chen, Zhiqing Zhong, Jie M. Zhang, Pinjia
671 He, and Ahmed E. Hassan. 2025. SWE-Eff: Re-
672 evaluating software AI agent system effectiveness
673 under resource constraints. *CoRR*, abs/2509.09853.

674 Shubham Gandhi, Manasi Patwardhan, Lovekesh Vig,
675 and Gautam Shroff. 2024. Budgetmlagent: A cost-
676 effective LLM multi-agent system for automating
677 machine learning tasks. In *Proceedings of the 4th*
678 *International Conference on AI-ML Systems, AIML-*
679 *Systems 2024*, pages 1–9.

680 Pengfei Gao and Chao Peng. 2026. More with less: An
681 empirical study of turn-control strategies for efficient
682 coding agents. In *Proceedings of the 48th Interna-*
683 *tional Conference on Software Engineering, ICSE*
684 *2026*.

685 Pengfei Gao, Zhao Tian, Xiangxin Meng, Xinchen
686 Wang, Ruida Hu, Yuanan Xiao, Yizhou Liu, Zhao
687 Zhang, Junjie Chen, Cuiyun Gao, Yun Lin, Yingfei
688 Xiong, Chao Peng, and Xia Liu. 2025. Trae agent:
689 An LLM-based agent for software engineering with
690 test-time scaling. *arXiv preprint arXiv:2507.23370*.

691 Tingxu Han, Zhenting Wang, Chunrong Fang, Shiyu
692 Zhao, Shiqing Ma, and Zhenyu Chen. 2025. Token-
693 budget-aware LLM reasoning. In *Findings of the As-*
694 *sociation for Computational Linguistics, ACL 2025*,
695 pages 24842–24855.

696 Sirui Hong, Mingchen Zhuge, Jonathan Chen, Xiawu
697 Zheng, Yuheng Cheng, Jinlin Wang, Ceyao Zhang,
698 Zili Wang, Steven Ka Shing Yau, Zijuan Lin, Liyang
699 Zhou, Chenyu Ran, Lingfeng Xiao, Chenglin Wu,
700 and Jürgen Schmidhuber. 2024. MetaGPT: Meta pro-
701 gramming for a multi-agent collaborative framework.
702 In *Proceedings of the Twelfth International Confer-*
703 *ence on Learning Representations, ICLR 2024*.

704 Nidhal Jegham, Marwan Abdelatti, Chan Young Koh,
705 Lassad Elmoubarki, and Abdeltawab Hendawi. 2025.
706 How hungry is AI? Benchmarking energy, water, and
707 carbon footprint of LLM inference. *arXiv preprint*
708 *arXiv:2505.09598*.

709 Huiqiang Jiang, Qianhui Wu, Chin-Yew Lin, Yuqing
710 Yang, and Lili Qiu. 2023. Llmlingua: Compressing
711 prompts for accelerated inference of large language
712 models. In *Proceedings of the 2023 Conference on*
713 *Empirical Methods in Natural Language Processing*,
714 pages 13358–13376.

715 Carlos E. Jimenez, John Yang, Alexander Wettig,
716 Shunyu Yao, Kexin Pei, Ofir Press, and Karthik
717 Narasimhan. 2024a. Swe-bench: Can language mod-
718 els resolve real-world github issues? In *The Twelfth*
719 *International Conference on Learning Representa-*
720 *tions, ICLR 2024*.

721 Carlos E. Jimenez, John Yang, Alexander Wettig,
722 Shunyu Yao, Kexin Pei, Ofir Press, and Karthik R.
723 Narasimhan. 2024b. Swe-bench: Can language mod-
724 els resolve real-world github issues? In *The Twelfth*
725 *International Conference on Learning Representa-*
726 *tions, ICLR 2024*.

727 Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio
728 Petroni, Vladimir Karpukhin, Naman Goyal, Hein-
729 rich Küttler, Mike Lewis, Wen-tau Yih, Tim Rock-
730 täschel, and 1 others. 2020. Retrieval-augmented
731 generation for knowledge-intensive nlp tasks. In *Ad-*
732 *vances in Neural Information Processing Systems*,
733 volume 33, pages 9459–9474.

734 Junwei Liu, Kaixin Wang, Yixuan Chen, Xin Peng,
735 Zhenpeng Chen, Lingming Zhang, and Yiling Lou.
736 2025. Large language model-based agents for soft-
737 ware engineering: A survey. *ACM Transactions on*
738 *Software Engineering and Methodology*.

739 Quinn McNemar. 1947. Note on the sampling error
740 of the difference between correlated proportions or
741 percentages. *Psychometrika*, 12(2):153–157.

742 OpenAI. 2025. Introducing swe-bench verified.

743 Joon Sung Park, Joseph C O’Brien, Carrie J Cai, and
744 1 others. 2023. Generative agents: Interactive simu-
745 lastra of human behavior. In *Proceedings of the 36th*
746 *Annual ACM Symposium on User Interface Software*
747 *and Technology*.

748 Shaolei Ren, Bill Tomlinson, Rebecca W Black, and
749 Andrew W Torrance. 2024. Reconciling the contrast-
750 ing narratives on the environmental impact of large
751 language models. *Scientific Reports*, 14(1):26310.

752 Gerard Salton and Christopher Buckley. 1988. Term-
753 weighting approaches in automatic text retrieval. *Information*
754 *processing & management*, 24(5):513–
755 523.

756 Noah Shinn, Federico Cassano, Ashwin Gopinath,
757 Karthik Narasimhan, and Shunyu Yao. 2023. Re-
758 flexion: Language agents with verbal reinforcement
759 learning. In *Advances in Neural Information Pro-*
760 *cessing Systems*, volume 36, pages 8634–8652.

761 StackOverflow. 2025. 2025 developer survey.

762 SWE-bench Team. 2025. SWE-bench: Benchmark
763 for real-world software engineering tasks. <https://www.swebench.com/>. Accessed: 2025.

765 Yicheng Tao, Yao Qin, and Yepang Liu. 2025. Retrieval-
766 augmented code generation: A survey with focus on
767 repository-level approaches. *CorR*, abs/2510.04905.

768 Chunqiu Steven Xia, Yinlin Deng, Soren Dunn, and
769 Lingming Zhang. 2025. Demystifying llm-based
770 software engineering agents. *Proceedings of the*
771 *ACM on Software Engineering*, 2(FSE):801–824.

772 Bruce Yang, Xinfeng He, Huan Gao, Yifan Cao, Xi-
773 aofan Li, and David Hsu. 2025. Codeagents: A
774 token-efficient framework for codified multi-agent
775 reasoning in llms. *arXiv preprint arXiv:2507.03254*.

776 John Yang, Carlos E Jimenez, Alexander Wettig, Kil-
777 ian Lieret, Shunyu Yao, Karthik Narasimhan, and
778 Ofir Press. 2024. Swe-agent: Agent-computer in-
779 terfaces enable automated software engineering. In
780 *Advances in Neural Information Processing Systems*,
781 volume 37, pages 50528–50652.

782 Qizheng Zhang, Michael Wornow, and Kunle Olukotun.
783 2025. Cost-efficient serving of llm agents via test-
784 time plan caching. *arXiv preprint arXiv:2506.14852*.

785 Wanjun Zhong and 1 others. 2024. Memorybank: En-
786 hancing large language models with long-term mem-
787 ory. In *Proceedings of the AAAI Conference on Arti-*
788 *ficial Intelligence*, volume 38.

789 A Experience Object Example

790 Figure 2 show an example experience object for the
791 django_django-16910 issue, constructed from the
792 execution of Mini-SWE-Agent with GPT-5-mini.

793 B Hyper-Parameter Tuning of GEAR

794 GEAR involves three hyper-parameters, including
795 τ_{sim} , $\tau_{\text{sel}}^{\text{gen}}$, $\tau_{\text{lower}}^{\text{sel}}$, and $\tau_{\text{upper}}^{\text{sel}}$). To tune them, we
796 randomly select 50 issues from the SWE-bench
797 dataset (Jimenez et al., 2024a) that are not included
798 in SWE-bench Verified or SWE-bench Lite to serve
799 as validation data.

800 **Tuning of τ_{sim} .** The similarity threshold τ_{sim} con-
801 trols the quality of experience retrieval in GEAR.

A threshold that is too low may retrieve irrelevant
experience objects, while a threshold that is too
high may retrieve too few, rendering GEAR inef-
fective due to insufficient guidance. To identify an
appropriate value, we conduct tuning experiments
on the validation data using Mini-SWE-Agent with
GPT-5-mini, chosen for its fast execution among
the evaluated agents. During this tuning, we ap-
ply experience injection without early termination.
The threshold is varied from 0.1 to 0.3 in steps of
0.05.

τ_{sim}	% Resolved	Total Cost
0.10	18%	\$1.46
0.15	24%	\$1.55
0.20	24%	\$1.56
0.25	20%	\$1.71
0.30	18%	\$1.99

Table 4: Resolution rate and total cost across different τ_{sim} settings.

Table 4 reports the resolution rate and total cost
for each threshold. As shown, $\tau_{\text{sim}} = 0.15$ achieves
both the highest resolution rate and the lowest cost,
and is thus adopted for all subsequent evaluations.
Tuning of τ^{gen} . The threshold τ^{gen} controls the
early termination of patch generation. Setting τ^{gen}
too high limits the effectiveness of early stopping,
whereas setting it too low may terminate genera-
tion prematurely and degrade task performance.
To identify an appropriate value, we conduct tun-
ing experiments on the validation data using Mini-
SWE-Agent with GPT-5-mini, chosen for its fast
execution among the evaluated agents. The thresh-
old is varied from 70 to 90 in increments of 5, and
we also evaluate the case without applying early
termination.

τ^{gen}	% Resolved	Total Cost
70	8%	\$1.24
75	8%	\$1.26
80	12%	\$1.36
85	20%	\$1.42
90	24%	\$1.50
w/o	24%	\$1.56

Table 5: Resolution rate and total cost across different τ^{gen} settings.

The results, summarized in Table 5, indicate that
a threshold of 90 achieves cost reduction while
maintaining the resolution rate comparable to that
of no early termination. Accordingly, we adopt
 $\tau^{\text{gen}} = 90$ for all subsequent experiments.

```
{
  "issue_id": "django_django-16910.traj",
  "task_description": "In Django 4.2, QuerySet.only() used with select_related() on a reverse OneToOneField fails to limit columns: all fields of the related model are still selected. Root cause: the ORM's only()/defer logic doesn't recognize or map the reverse OneToOne select_related join correctly (relation/alias handling), so column trimming isn't applied during query construction.",
  "execution_summary": [
    "Find Files: Locate QuerySet.only implementation",
    "Read Code: Inspect QuerySet.only method in django/db/models/query.py",
    "Read Code: Inspect add_immediate_loading and related deferred/immediate logic in django/db/models/sql/query.py",
    "Analyze Logic: Examine get_select_mask and _get_only_select_mask to find how select masks are built",
    "Modify Code: Apply initial patch to _get_only_select_mask to use related field when opts.get_field returns a reverse accessor",
    "Run Tests: Compile-modified module to ensure syntactic correctness (py_compile)",
    "Run Tests/Debug Issue: Reproduce query generation with a small script and inspect resulting errors",
    "Modify Code: Replace _get_only_select_mask with a defer-mirroring implementation handling forward and reverse relation masks; verify change via git diff"
  ],
  "evaluation_result": "Pass",
  "confidence": 90,
  "confidence_reason": "High confidence (90%) stems from: (1) Root cause identified: ORM's only()/defer logic fails to map reverse OneToOne select_related joins during query construction, causing column trimming to be skipped. (2) Solution mirrors proven defer() implementation pattern, ensuring consistency with forward relations. (3) Fix addresses the core issue at _get_only_select_mask level by properly handling reverse accessors via related field resolution. (4) Systematic approach: code inspection to logic analysis to targeted modification to verification. Key insight: reverse relations require explicit alias/relation mapping that forward relations handle automatically.",
  "created_at": "2025-11-09T23:55:25.286001",
  "metadata": {
    "original_count": 8,
    "summarization_method": "llm",
    "model": "openai/gpt-5-mini"
  }
}
```

Figure 2: Experience object of the django_django-16910 issue.

Tuning of $\tau_{\text{lower}}^{\text{sel}}$ and $\tau_{\text{upper}}^{\text{sel}}$. The two thresholds control early termination during patch selection. Specifically, $\tau_{\text{upper}}^{\text{sel}}$ defines the confidence boundary above which a patch is assumed correct and no further generation is needed, while $\tau_{\text{lower}}^{\text{sel}}$ defines the boundary below which the issue is considered too difficult to resolve completely.

Unlike the previous thresholds, tuning of these two is conducted on the validation set using Trae Agent with GPT-5-mini, since Mini-SWE-Agent does not include a patch selection phase. For $\tau_{\text{lower}}^{\text{sel}}$, we vary its value from 0 to 50 in increments of 10, leaving $\tau_{\text{upper}}^{\text{sel}}$ unset during this tuning. For $\tau_{\text{upper}}^{\text{sel}}$, we vary it from 80 to 100 in increments of 5, leaving $\tau_{\text{lower}}^{\text{sel}}$ unset.

Tables 6 and 7 summarize the results. Based on these experiments, we select $\tau_{\text{lower}}^{\text{sel}} = 40$, which achieves the highest resolution rate at the lowest cost, and $\tau_{\text{upper}}^{\text{sel}} = 90$, which similarly balances

	Confidence	Solve Rate	Cost
0	30%	\$15.8	
10	30%	\$15.2	
20	30%	\$14.7	
30	30%	\$13.4	
40	30%	\$12.8	
50	26%	\$12.4	

Table 6: Resolution rate and total cost across different $\tau_{\text{lower}}^{\text{sel}}$ settings.

maximal resolution with minimal cost.

In patch selection part, the higher threshold determines the boundary where the patch is assumed to be correct and needs no regeneration, and the lower threshold determines the boundary where the issue is too hard to be complete. To determine the thresholds, we use the same subset of SWE-bench and do experiment on Trae-agent with GPT-5-mini as backend. According to Table 6 and 7, we choose 40

834
835
836
837
838
839
840

841
842
843
844
845
846
847
848

849
850
851
852

853
854
855
856
857
858
859
860
861

Confidence	Solve Rate	Cost
100	30%	\$15.8
95	30%	\$15.0
90	30%	\$14.3
85	24%	\$12.3
80	20%	\$10.8

Table 7: Resolution rate and total cost across different $\tau_{\text{upper}}^{\text{sel}}$ settings.

as the lower threshold and 90 as the higher threshold, for they are the thresholds that minimize cost without compromising performance.

C Detailed Comparison with Baselines

Table 8 presents the effects of GEAR and the two baseline methods on task performance and cost across different SE agents.

Method	% Resolved	# API Calls	# Input Tokens	# Output Tokens	Total Cost
Agentless + GPT-5-mini					
Naive-GEAR	+1.0%	-24.0%	-48.7%	-49.9%	-53.0%
GEAR	+7.8%	-26.4%	-51.8%	-51.0%	-55.1%
Agentless + DeepSeek-V3.2					
Naive-GEAR	+1.2%	-21.8%	-26.3%	-28.7%	-25.8%
GEAR	+7.2%	-25.5%	-31.9%	-35.0%	-31.8%
Mini-SWE-Agent + GPT-5-mini					
Turn-control	-13.0%	-33.5%	-60.1%	-38.1%	-50.5%
Naive-GEAR	-0.6%	+2.7%	+16.0%	+17.0%	-1.1%
GEAR	+1.0%	-7.9%	-13.7%	-3.7%	-19.4%
Mini-SWE-Agent + DeepSeek-V3.2					
Turn-control	-11.0%	-28.4%	-46.8%	-17.1%	-41.6%
Naive-GEAR	-4.2%	-4.2%	-8.5%	-4.0%	-8.2%
GEAR	+0.6%	-8.4%	-13.6%	-4.4%	-19.3%
Trae Agent + GPT-5-mini					
Turn-control	-10.0%	-34.6%	-34.9%	-32.5%	-34.6%
Naive-GEAR	-9.0%	-34.1%	-34.7%	-32.1%	-31.7%
GEAR	0.0%	-29.9%	-30.4%	-28.0%	-28.2%
Trae Agent + DeepSeek-V3.2					
Turn-control	-8.8%	-39.0%	-39.5%	-39.5%	-38.7%
Naive-GEAR	-7.0%	-18.2%	-34.0%	-27.0%	-36.4%
GEAR	-0.2%	-26.5%	-37.7%	-28.2%	-36.7%
Average					
Turn-control	-10.7%	-33.9%	-45.3%	-31.8%	-41.4%
Naive-GEAR	-3.1%	-16.6%	-22.7%	-20.8%	-26.0%
GEAR	+2.7%	-20.8%	-29.9%	-25.1%	-31.8%

Table 8: Impact of GEAR and baseline methods on SE agents' task performance and cost.